

Polymarket Microstructure and Trading

Notes on manual trading, liquidity, and inefficiencies in sports prediction markets.

Ryan Wu

May 2026

Disclaimer: *These notes are intended solely for educational and informational purposes and do not constitute financial advice, investment advice, or a recommendation to trade prediction markets. Trading prediction market contracts involves substantial risk, including the possible loss of capital.*

Introduction

From November 2024 to October 2025, I analyzed Polymarket's sports markets, and the following notes summarize observations and strategies derived from the analysis.

It is important to note that the following strategies are written from the perspective of a **manual trader**, a trader who does not automate the process of placing or canceling orders. Unfortunately, the strategies that I will discuss are no longer as profitable today as they were during my period of analysis. However, there are non-trivial insights that new traders can use to formulate their own strategies that adapt to today's more competitive prediction market landscape.

The category of markets that I found to be most profitable to trade was sports. My reasons are as follows:

1. Sports markets tend to have higher volume than other categories of markets, and higher volume markets generally have more frequent mispricings.
2. There exist many more bookmakers and alternative prediction markets offering sports odds, which provide reference points for estimating fair value in sports markets.
3. Sports games represent periods of planned volatility, while other market categories have unexpected periods of volatility.
4. There are more clearly defined mathematical relationships between different sports markets than in other markets.

With that, we are ready to begin.

1. Overview and Definitions

I will need to introduce several terms that I developed to help describe market structure and trading strategies for sports markets. Because my terminology is nonstandard, understanding the definitions will be important for the remainder of the text.

I define a **primary market** as a market that can be resolved using the outcome of one single game or event. In basketball, that would correspond to a single game. In tennis, a single match. In Formula 1, a single race, although there are nuances to this definition. At the time of writing, primary markets are generally found under the tab on Polymarket labeled "Games."

Primary markets are populated by **market makers**. I define a market maker as a trader who automates the process of placing orders, usually very quickly and accurately around fair value, even during periods of volatility.

The speed at which you are able to place orders on Polymarket is significant because open orders at the same price are filled in the order they were placed (first in first out). Market makers usually place open orders of absurdly large size, so even placing identical orders to market makers will likely not create any trading volume for you.

Furthermore, highly liquid primary market contracts usually have a spread of $1c$, so even if you could get your orders filled, you would only earn an expected profit of $0.5c$ /contract, assuming the true fair value is symmetrically distributed within the spread. Relative to a contract trading near fair value at $50c$, that would only yield an edge of 1%, which underperformed other strategies¹.

Because I was unwilling and unable to compete with faster and more automated market makers for small profits, I searched for other opportunities.

My trading strategy revolved almost exclusively around **derivative markets**. I define a derivative market as a market whose value depends on the outcome of multiple games or events. For example, the “NBA Champion” market derives its value from the outcomes of almost every NBA game.

It is important to note that derivative markets can become functionally equivalent to primary markets. For example, during Game 7 of the NBA Finals in 2025, the “NBA Champion” market is functionally equivalent to the primary market “Thunder vs Pacers.”

My definitions for primary market and derivative market are not perfect, and there are ambiguous cases. For example, the “Miami Grand Prix Winner” market can be thought of as dependent on the “Miami Grand Prix Qualifying Winner” market, meaning it would only behave like a primary market at the conclusion of qualifying.

The exact definitions are less important than understanding the structural characteristics that make derivative markets more attractive to trade.

2. Hedging

It may seem counterintuitive to discuss hedging before trading strategies, but I believe the unique characteristics of prediction market contracts make such an ordering very necessary.

Compared with traditional stock or option contracts, prediction market tokens are extremely high-volatility, short-term assets. You should approach prediction markets with a clear understanding of your personal risk tolerance.

¹ This does not take market rewards into account, which reward liquidity providers for placing orders near the midpoint of the spread.

I would even recommend that manual traders split their bankroll across multiple accounts, where each account contains an amount of money they are willing to risk entirely on a single trade. This provides an additional benefit because you can use the “Max” option to buy contracts immediately without spending time calculating position sizing during periods of volatility². The second of computation time that this saves is not negligible, as will become evident later on.

In most cases, we will use liquid primary markets to hedge. From our earlier discussion regarding market maker profit, it should already be clear that hedging through liquid primary markets generally requires paying a premium of approximately $0.5c$ /contract when the spread is $1c$.

Initially, I naively believed that this premium made hedging strictly -EV. If we consider only the raw dollar value of money, this would indeed be true. Purchasing a contract at fair value has expected value 0, so paying $0.5c$ /contract above fair value is necessarily a loss in expectation.

However, this reasoning ignores the future earning power of capital.

Even today, at account sizes below approximately \$50,000, an annual return of 20% is, in my opinion, a conservative estimate for what a strong manual trader can achieve³. Using this assumption, we can construct a simple example showing why hedging may still be profitable despite paying above fair value.

Let

- A represent current account size,
- P represent profit if the trade wins,
- L represent loss if the trade loses,
- and p represent the probability that the trade wins.

Assume that the original market mispricing has already disappeared, so the market prices are now fair. Assume also that we can liquidate on both sides at the fair price. Under these assumptions,

$$EV = Pp - L(1 - p) = 0$$

Now consider future earning power.

If we risk a loss of L , we are not only risking the nominal amount L , but also the future earning potential associated with that capital. Denote this future earning value by L_f . Similarly, if we profit P , we also gain the future earning potential associated with that profit, which we denote by P_f .

² At the time of writing, Polymarket does not seem to have the Max option for buying shares anymore.

³ The estimate is entirely based on my own analysis of Polymarket sports markets today, informed by my trading experience.

Define the future earning multipliers c_L and c_P such that

$$L_f = c_L L$$

and

$$P_f = c_P P$$

We can now define an adjusted expected value EV_f that incorporates future earning power:

$$EV_f = (P + P_f)p - (L + L_f)(1 - p)$$

Substituting gives

$$\begin{aligned} EV_f &= (P + c_P P)p - (L + c_L L)(1 - p) \\ &= (Pp - L(1 - p)) + c_P Pp - c_L L(1 - p) \end{aligned}$$

Since

$$Pp - L(1 - p) = 0$$

we obtain

$$EV_f = c_P Pp - c_L L(1 - p)$$

If future earning power were independent of account size, then we would have

$$c_P = c_L = c$$

which would imply

$$EV_f = cPp - cL(1 - p) = c(Pp - L(1 - p)) = 0$$

However, in practice, the future earning power of money is negatively correlated with account size. Losing capital generally increases the marginal value of each remaining dollar, while gaining capital decreases it. Therefore,

$$c_L > c_P$$

Substituting this inequality gives

$$\begin{aligned} EV_f &= c_P Pp - c_L L(1 - p) \\ &= (c_P - c_L)Pp + c_L (Pp - L(1 - p)) \\ &= (c_P - c_L)Pp < 0 \end{aligned}$$

Thus, even when a hedge appears to be negative EV in nominal terms, hedging may still be profitable once future earning power is considered. In particular, if the cost of hedging is less than

$$(c_L - c_P)Pp$$

then hedging increases expected future value.

The above framework relies on many assumptions, so in practice it is better to develop rules of thumb based on personal risk tolerance rather than applying the formulas mechanically. The important intuition is that future earning power is not linearly related to account size. For example, the marginal value of capital changes far more going from $A = 100$ to $A = 200$ than from $A = 100,000$ to $A = 100,100$.

As a result, you should generally be much more willing to hedge when a loss would represent a significant fraction of your bankroll.

Finally, from personal experience, the value of peace of mind truly cannot be discounted.

Aside

In an interview, I was asked something along the lines of the following question:

There is a match being played between Team A and Team B. The market's pricing is 70% for Team A winning and 30% for Team B winning. Your friend is a fan of Team B. He offers you 10,000 Team A shares at 50%. A stranger offers you 400,000 Team B shares at 25%. You must decide which offer to take. You cannot take both.


You can provide a perfectly acceptable answer using adverse selection to justify why your friend's offer is better, but let us assume that we are entirely confident in the market's price being the fair value. If this were the case, it is obvious that the stranger's offer returns a higher raw dollar EV. However, the better offer to take depends also on the future earning potential of our money, which differs from person to person. For example, if I only had \$100,000, I would prefer my friend's offer. If I had \$100M, I would take the stranger's offer.

We are now ready to begin discussing trading strategies.

3. Single Elimination Tournaments (SETs)

Some events I traded in this category include the NBA playoffs, NFL playoffs, Tennis Grand Slam tournaments, and March Madness. Single elimination tournaments are one of the most common structures in sports, so many derivative markets fall into this category. Common examples include markets such as "Who will win the Finals?," "Who will win the Eastern Conference?," and "Will a #1 seed win the Finals?"

We will use the following tournament structure as an example:



Assume A vs B is the first semifinal match and C vs D does not begin until after A vs B concludes.

During the live semifinal match between A and B , the Finals market for A and B behaves approximately like a leveraged version of the semifinal market. Intuitively, the probability that A or B wins the tournament changes primarily through the change in probabilities that A and B win the semifinal round.

We can formalize this relationship as follows.

Let T_s and T_f represent the events that team T wins the semifinal and Finals respectively for $T \in \{A, B, C, D\}$.

Let p_{Ts} and p_{Tf} represent the pregame probabilities that team T wins the semifinal and Finals respectively.

As close to the start of the A vs B game as possible, but before the game begins, store the value of p_{Ts} and p_{Tf} for $T \in \{A, B\}$.

During the game, we approximate the Finals probability using $P(T_f) \approx \frac{P(T_s)}{p_{Ts}} p_{Tf}$ for $T \in \{A, B\}$.

Note that we calculate $P(T_s)$ using the midpoint of the spread in the A vs B semifinal primary market.

Let's derive why this approximation works well.

We know

$$P(A_f) = P(A_s) \cdot P(A_f | A_s)$$

Expanding the conditional outcomes for the opposite side of the bracket gives

$$P(A_f) = P(A_s)[P(A_f | A_s, C_s) + P(A_f | A_s, D_s)]$$

This equation should make intuitive sense. The probability that A wins the tournament equals the probability that A reaches the Finals multiplied by the probability that A wins the Finals conditional on which team it faces.

If we assume that

$$P(A_f | A_s, C_s) + P(A_f | A_s, D_s)$$

remains constant during the live semifinal match, we can rewrite the equation as

$$P(A_f) = P(A_s) \cdot c$$

for some constant c .

In the moment we stored the values for p_{Af} and p_{As} , it was true that $P(A_f) = p_{Af}$ and $P(A_s) = p_{As}$, so

$$p_{Af} = p_{As} \cdot c \Rightarrow c = \frac{p_{Af}}{p_{As}}$$

Substituting back gives

$$P(A_f) = P(A_s) \cdot \frac{p_{Af}}{p_{As}} = \frac{P(A_s)}{p_{As}} p_{Af}$$

which yields the approximation

$$P(T_f) \approx \frac{P(T_s)}{p_{Ts}} p_{Tf}$$

which matches our expression from earlier.

The assumption that

$$P(A_f | A_s, C_s) + P(A_f | A_s, D_s)$$

does not change significantly during A vs B is not unreasonable. In the vast majority of cases, there should not be any significant event during the match that drastically changes the bookmakers' odds for A in the Finals. However, there certainly are events that can alter these probabilities.

For example, suppose A wins, but A 's star player sustains a season-ending injury at the conclusion of the semifinal. In that case, the probability that A wins the Finals would decrease relative to the case where the injury did not occur. I call events such as these, events that significantly alter the probability of a derivative market without significantly altering the probability of the ongoing primary market, **soft volatility events**.

There are effectively an infinite number of possible soft volatility events, although each individual event generally occurs with low probability.

As a manual trader, I handled soft volatility events using my personal approximations and intuition. You may be able to outperform my approximations using a model, but it will likely not improve your profitability significantly because the additional mispricings you find will be at the fringe of the fair value anyway. Furthermore, obtaining meaningful training data for the effect of an effectively infinite set of soft volatility events would be difficult.

Taking the strategy one step further, we can also use the result of A vs B to trade C and D in the Finals market. Consider the case where B is a heavy favorite to win both the semifinal and to win the tournament. If A were to upset B in the semifinals, this would intuitively increase the probability that either C or D wins the tournament. Assuming there are no soft volatility events, we can derive the exact changes in the Finals probabilities for C and D based entirely on the live probability of the semifinal A vs B .

From the previous derivation, we already have four equations of the form

$$\begin{aligned} P(A_f) &= P(A_s)[P(A_f | A_s, C_s) + P(A_f | A_s, D_s)] \\ P(B_f) &= P(B_s)[P(B_f | B_s, C_s) + P(B_f | B_s, D_s)] \\ P(C_f) &= P(C_s)[P(C_f | C_s, A_s) + P(C_f | C_s, B_s)] \\ P(D_f) &= P(D_s)[P(D_f | D_s, A_s) + P(D_f | D_s, B_s)] \end{aligned}$$

At all times, we know $P(T_s)$ for $T \in \{A, B, C, D\}$ using the live primary market probabilities for the semifinal matches.

Notice that

$$P(A_f | A_s, C_s) = 1 - P(C_f | C_s, A_s)$$

with analogous identities for the remaining conditional probabilities.

This gives a system of 12 equations and 12 variables, allowing the remaining algebra to be completed and exact update formulas to be derived for the Finals probabilities of C and D using only the pregame Finals probabilities and the live semifinal probability of A vs B under the assumption that no soft volatility events occur.

These strategies worked because some derivative markets did not have active market makers, and those that did often operated with a delay of a few seconds relative to the market makers in the primary markets. This created opportunities to purchase mispriced orders in derivative markets before they repriced.

Today, many traders have automated this process or reduced their information delay to nearly zero relative to the primary markets. As a result, filling mispriced derivative market orders has become substantially more difficult, especially when using standard streaming channels.

4. Limit Orders

While SETs were usually my preferred method of trading, there are periods during the year where there are not enough high-volume sporting tournaments available to trade. During these periods, you need to use limit orders to avoid leaving capital idle.

For a manual trader, limit orders are generally safest to place during low-volatility periods because you do not receive information as quickly as market makers. Placing limit orders during a live game exposes you to the risk that your own open orders become the mispriced orders discussed in the previous section.

With that being said, let us discuss how to place profitable limit orders.

Which Markets Do I Choose?

This question becomes easier to answer once we understand our actual objective when placing limit orders.

Very roughly speaking, our goal is to maximize

$$(number\ of\ fills) \cdot (fill\ price - fair\ value)$$

while simultaneously minimizing the probability that our bid falls below fair value and minimizing the amount of capital tied up in open orders.

We will address each of these considerations throughout the remainder of this section.

There is usually a negative correlation between the number of fills you receive and the amount of edge you earn per fill. As you decrease your bid farther away from fair value, your expected edge per fill increases, but the expected number of fills decreases. Furthermore, once you decrease your bid below the highest bid of size from the existing open orders, your expected fills will drop very sharply⁴.

Because markets with tighter spreads also tend to have higher trading volume, this creates a broader tradeoff between trading volume and spread size as well. Thus, the ideal market for placing limit orders is one with enough trading volume to regularly fill orders, but also a sufficiently large spread to generate a meaningful edge per fill.

Markets with extremely low volume and very wide spreads are usually not worth trading because fills occur too infrequently. Conversely, highly liquid markets such as primary markets or very large derivative markets like “Who will win the NBA Championship?” or “Who will win the Super Bowl?” often have spreads that are too small to offer meaningful profit potential for manual traders.

⁴ We only discuss bids for simplicity. Because prediction market contracts are binary, every bid on one side of the market is equivalent to an ask on the complementary side.

This leaves moderately liquid derivative markets as the most attractive candidates for profitable limit order trading.

Markets with Unusually High Fill Rates

The most profitable markets I traded using limit orders were usually markets with many low-probability outcomes, specifically when buying No shares. One example is the “Club World Cup Top Goalscorer” market. An interesting pattern I observed in markets such as these was that No token buyers often appeared to achieve significantly better fills relative to fair value than Yes token buyers.

There are several possible explanations for this. It may simply be that, even after accounting for the Polymarket spread, the implied odds are still more attractive than traditional gambling alternatives. It may also be the case that traders looking to gamble are more attracted to low-probability Yes tokens than high-probability No tokens. However, that is a question for behavioral psychologists.

These markets also provide an important structural advantage when holding No tokens: protection against volatility spikes. Consider the “Club World Cup Top Goalscorer” market again. Outside of the planned volatility during matches, periods during which I would generally cancel all open orders, the largest volatility events usually came from injuries or other events that effectively removed a player from contention.

If you instead placed open orders to buy Yes tokens on individual players, you would expose yourself to substantial downside risk from these volatility events. In practice, there is generally much more room for a low-probability Yes token to collapse in value than to rapidly increase.

As a result, unless you are actively monitoring news events, it is extremely important to consider how your open orders may be exposed to unplanned volatility.

Finding Fair Odds

We now understand, at least on an abstract level, that the best markets for placing limit orders generally have a reasonable amount of volume, reasonably large spreads, and preferably many low-probability options.

The next step is determining a fair value from publicly available information so that we can establish a reference price for our limit orders.

When the fair value of a derivative market can be derived almost entirely from primary market probabilities, aggregation is relatively straightforward. In these cases, Polymarket and Kalshi are usually sufficient because both platforms generally maintain high liquidity in primary markets.

However, for derivative markets whose probabilities cannot be derived directly from primary markets, such as “Top Goalscorer,” “MVP,” or “Formula 1 H2H,” we must instead rely on sportsbook odds. Fortunately, most reasonably active sports derivative markets have equivalent or near-equivalent markets

available at one or more sportsbooks. I suspect this is partially because Polymarket would struggle to attract liquidity in a derivative market without sportsbooks first establishing an approximate reference price.

Sportsbook odds almost always exceed fair value because sportsbooks need to earn a profit. I define **implied probability** as the probability obtained by converting sportsbook odds into probability space. As a result, when summing the implied probabilities of all outcomes in a market, the total should always exceed 100%⁵.

For high-volume primary markets with only a few outcomes, this total may be approximately 105% – 110%. However, for low volume derivative markets with many low-probability options, especially when few sportsbooks offer equivalent markets, the total implied probability will be larger, perhaps exceeding 150%⁶. Clearly, we therefore need to produce normalized probabilities, probabilities adjusted so that the outcomes sum to exactly 100%.

The standard normalization procedure is simple: divide each implied probability by the sum of all implied probabilities. This usually provides a good approximation in markets where the probabilities of the different outcomes are relatively similar. For that reason, this method often performs well for liquid primary markets of competitive games.

However, this approach does not fully account for sportsbooks' tendency to aggressively overprice low-probability outcomes. An option with a true probability near 50% may have an implied probability near 55%, while a 1% true probability outcome in the same market may have an implied probability closer to 2%. Thus, applying the same normalization scaling factor uniformly across all outcomes is not always appropriate.

I could not give you an exact formula for converting sportsbook odds into fair probabilities. In practice, I often adjusted probabilities heuristically until the resulting distribution appeared approximately reasonable to me. Fortunately, having an extremely precise estimate of fair value is not always necessary. In many cases, limit orders to buy can be placed sufficiently below estimated fair value anyway.

I am certain many firms and traders have already analyzed this privately, but one potentially valuable dataset would consist of sportsbook odds as input features and corresponding Polymarket market midpoints as labels. I would use such a dataset to train a pricing model for derivative markets.

Further Market Selection

I would like to briefly describe two more relevant concepts: **surprise factor** and **hype factor**.

⁵ For the remainder of this section, assume the market outcomes are mutually exclusive and collectively exhaustive, so that exactly one outcome resolves Yes.

⁶ This is sometimes caused by the specific payout rules of derivative markets in casinos that allow for multiple winners, but we will cover this in a later section.

I would like to preface that surprise factor is a term that I invented to describe how subjectively surprising some event feels to me. However, I think the concept generalizes reasonably well because people's intuitions do not differ very dramatically. Similarly, I think hype factor generalizes well because people consume similar media.

Here is an example of surprise factor. At the time of writing, the 2026 World Cup is coming up soon. If I were to ask you how surprised you would be if France won the World Cup, you would likely say that you would not be very surprised at all, despite France's pre-tournament odds sitting near 18%. Yet, if I asked you how surprised you would be if an 18% underdog won some individual match, you would likely say that you would be very surprised.

I believe surprise factor is influenced not only by probability itself, but also by the number of alternative possibilities a person considers. Because France is currently one of the favorites to win the World Cup, all the other low-probability countries collectively make France winning feel more likely. However, in a single match, the alternative to the underdog winning is the 82% favorite, so the favorite feels much more dominant psychologically.

Hype factor is more straightforward and likely easier to measure, perhaps using social media mentions. For example, if I were to ask who the most hyped player in the 2026 World Cup is, many people would probably answer Messi.

I try not to let my own surprise factor or hype factor influence my estimate of fair value, and I would strongly recommend you do the same. However, understanding that many traders do allow these factors to influence them should give you a good idea of which orders are most likely to be filled.

Because manual trading carries significant time costs, you can prioritize market options with low surprise factor and high hype factor when placing orders.

Placing the Order

You will ultimately need practical experience placing limit orders to gauge how aggressively above or below the current midpoint you can reasonably price while still getting filled.

There will often be automated orders in derivative markets, so one simple approach is to place orders slightly better than the first open order of significant size already in the order book. You could also use a fixed spread relative to your estimated fair value.

It is best to experiment and determine which approach you personally prefer.

5. Liquidation

Assuming the previous strategies have worked, you should now hold contracts rather than cash.

Time

The expected value of a position is approximately

$$(number\ of\ contracts) \cdot (fair\ value)$$

So if we can liquidate above fair value, we should certainly do so.

When liquidation price = fair value, liquidation can still be desirable because we continue to face the same future earning power considerations discussed in Section 2. Holding the position exposes us to downside risk that may reduce future profit potential more than the upside increases it.

Even when liquidation price < fair value, liquidation may still be preferable because it frees both capital and time. Unlike hedging, liquidation immediately releases capital to be redeployed elsewhere rather than requiring additional capital to establish a hedge. As a result, we should generally be willing to liquidate at lower EV prices than we would accept when hedging.

Let's now look at some actionable liquidation examples.

Example 1

In the previous section, I discussed buying No tokens in markets with many low-probability options. While this can be a profitable strategy, one downside is that No positions often tie up large amounts of capital while remaining difficult to hedge or liquidate efficiently.

This is especially true in markets such as “Top Goalscorer,” “MVP,” and similar derivative markets where perfect hedges are usually impossible. One possible solution is to provide liquidity across multiple No tokens in a single winner-take-all market and use Polymarket’s Convert and Merge functions to partially liquidate positions.

For example, consider a market with A , B , ..., J . Suppose we hold 500 shares of A No and 300 shares of B No. We can convert 300 shares of B No into Yes shares on all remaining options. This produces 300 shares of A Yes, which can then be merged against 300 shares of A No, immediately freeing \$300 in capital.

Example 2

Once the outcome of a market becomes effectively certain, the winning token will often trade near 99.9c. Many traders, including myself, are willing to sell at these prices because the several hours required for formal market resolution may be worth more to them than the remaining 0.1c/contract.

This naturally raises the question of why other traders are still willing to buy at 99.9c.

The answer is that the value of time depends heavily on the individual trader. In general, the more capital a trader has, the less likely they are to sacrifice EV simply to recover liquidity slightly earlier.

Example 3

Liquidation also protects against positions that become too large to hedge comfortably over time.

For example, suppose you purchase a Finals derivative market token early in the tournament at 5c using 2% of your account value. Initially, hedging the position is relatively manageable. However, suppose the underlying event evolves favorably and the position eventually represents 20% of your account value. At that point, fully hedging the position may itself require committing another 20% of your account.

Low-probability positions can gradually evolve into extremely large exposures. Unless you were originally prepared to let the full position ride from the beginning, continuously hedging instead of liquidating may eventually force you into taking substantially more risk than you intended.

5. Technicalities

This section is about finding unique ways to identify mispricings.

Resolution Rules

Resolution rules are one of the most important things to analyze before trading any market. Polymarket and sportsbooks often resolve what appears to be the “same” market in very different ways. Consider top goalscorer markets, for example. I have seen cases where the tiebreaker was alphabetical order of listed last name, and other cases where the tiebreaker was assists followed by alphabetical order.

Furthermore, many Polymarket probabilities are implicitly derived from sportsbook odds in “equivalent” markets. However, sportsbooks are often not actually offering the same market structure that Polymarket is offering. This is because sportsbooks often use dead heat settlement for ties rather than explicit tiebreakers⁷. This can create significant discrepancies between sportsbook odds and Polymarket probabilities.

As a simple example, consider a Champions League top goalscorer market using alphabetical order of listed last name as the tiebreaker. Haaland could be tied with Dembele for the most goals scored, but if Manchester City had already been eliminated from the tournament, sportsbooks could still continue offering odds on Haaland while Polymarket would have already resolved Haaland to No. This is a type of edge case that you must understand before entering a position.

Many traders believe edge cases are unimportant because they occur infrequently. However, the EV of understanding edge cases is undervalued. You are rarely going to make significant money from obvious

⁷ Dead heat settlement is a sportsbook resolution method in which tied winners split payouts proportionally instead of resolving a single winner through tiebreakers.

outcomes and straightforward rules. But if you correctly prepare for an edge case before it occurs, the payoff can be extremely large.

Negative Risk and “Other”

Most derivative markets are **winner-take-all** markets, meaning exactly one outcome resolves positively. These markets often support **negative risk**, meaning No shares can be converted into equivalent Yes shares on all other outcomes.

However, you must be very careful about whether a market supports **complete negative risk**. A market has complete negative risk if future outcomes can still be added to the market after trading begins. In these markets, converting No shares into Yes shares earlier means you would also receive shares in newly added outcomes later. The presidential election market is an example of a market with complete negative risk.

This creates an important implication: if the Yes shares in a complete negative risk market do not sum to 100c, this may simply mean that additional outcomes have not yet been added to the market⁸.

The fundamental distinction between markets with complete negative risk and markets without complete negative risk is often the presence of an “Other” option. Markets without complete negative risk frequently include an Other option to account for unlisted possibilities. In contrast, it may seem like markets with complete negative risk have no need for an Other option because any other possibilities can simply be added.

However, that is not the case. Consider the 2024 FIDE World Blitz Championship⁹.

The 2024 FIDE World Blitz Championship market was a winner-take-all market with complete negative risk. However, Magnus Carlsen and Ian Nepomniachtchi agreed to share the title, resulting in the Other option winning.

You may believe the “fair” outcome would have been for both Magnus Carlsen and Ian Nepomniachtchi to resolve Yes. However, understanding beforehand that the market structure required exactly one outcome to resolve to Yes could potentially have been very profitable.

It is also important to carefully read the exact resolution language for Other tokens. For example, many markets contain clauses similar to:

“If this event is cancelled, or not completed by March 31, 2025, 11:59 PM ET, this market will resolve to No.”

In other markets, cancellation may instead produce a 50-50 resolution.

⁸ It may also partially reflect the time value of money. A market resolving far in the future may have prices lower than the true probability should suggest.

⁹ <https://polymarket.com/event/fide-world-blitz-champion>

These scenarios are not negligible edge cases. Real-world events are much messier than the clean structure implied by prediction market contracts. Situations like these have happened before and will continue to happen in the future, so it is important to already have a plan for how to react when they do occur.

6. Final Remarks

Everything above is conceptually straightforward, and I am intentionally choosing to leave out highly specific strategies for individual markets. The purpose of these notes is to outline the core tools needed to operate in prediction markets. Once you understand which markets derive from others, how edge cases affect resolution, and how to manage risk, you should be capable of developing profitable strategies on your own.

The days of profitable manual trading are likely coming to an end, but I do not believe what I have written is obsolete. Entering prediction markets with trading tools but without practical experience is only marginally better than gambling. Learn carefully, test cautiously, and stay safe. Ultimately, the most valuable thing you gain from prediction markets may not be the money, but the education that comes from operating in a real financial market.